

Quadratic Sieve Algorithm

a.k.a. the Second Fastest Integer Factoring Algorithm in the West
(and trying to get it to work)

Nils Olsson

Spring Semester, 2021

Contents

1	Introduction	2
2	Factoring via Congruences of Squares	3
2.1	Trial Division	3
2.2	Fermat's Method	4
2.3	Congruences of Squares	6
2.4	Dixon's Random Squares Algorithm	7
2.5	The Quadratic Sieve Algorithm	8
3	Implementation	10
3.1	Implementing in Rust	10

Abstract

This report serves as a summary of the history and mathematical theory behind square factoring functions, as an loose introduction to the Rust programming language, and as a retrospective on my own implementing of the quadratic sieve (QS) algorithm.

Chapter 1

Introduction

The factorization of integers has been an especially well studied problem, and not just for centuries but in fact millennia. Greek mathematicians studied the problem, even proving the fundamental theorem of arithmetic: that every integer has a unique prime factorization. As a corollary the Greeks had also extensively studied the prime integers, with the sieve of Eratosthenes being one of the most important algorithms for generating primes ever invented. The so called father of geometry Euclid of Alexandria (whose now self-titled “Euclid’s lemma” was foundational for proving integer factorization uniqueness) described in his *Elements* the Euclidean algorithm for calculating the greatest common divisor of two integers. These ancient algorithms, and many others, comprise the fundamental building blocks for which all modern integer factorization algorithms are made.

Chapter 2

Factoring via Congruences of Squares

In this section we will cover several factorization algorithms and the fundamental math that they are based upon. These individual topics lead-up to covering the quadratic sieve (QS) algorithm in detail.

Something to pay special attention to is the conditions on the parameters for various algorithms, because many of the factoring algorithms which we will discuss are *not* complete factoring algorithms. In fact, all but one is: trial division.

2.1 Trial Division

At its simplest, the task of fully finding the prime factorization of integer, and the task determining if said integer is prime is the same task. The simplest method of doing so is what we call *trial division*. If n is the integer we wish to factor, then with trial division we consider all the integers from 2 to the first integer greater-than or equal-to the square root of n (since \sqrt{n} is precisely the largest number which divides n). The process is as follows:

With $d = 1, 2, 3, \dots$ determine if n is divisible by d . If n is divisible by d , we take note of both d and e the power of d which divides n , and let $n \leftarrow n/d^e$. We only stop when $d \not\leq \lceil \sqrt{n} \rceil$.

With the algorithm having terminated, we know the prime power factorization of n given the d divisors and their corresponding e exponents that we collected. If none of the d divisors ever divided n then we can say with absolute certainty that n is prime. This algorithm is simple enough to implement in Python in but a few lines of code:

```
# Factor an integer via trial division
```

```

# @param n Integer to factor
# @return Its prime power factorization
def factor_trial_division(n):
    fs = []
    d = 2
    root = ceil(sqrt(n))
    while d ≤ root and n > 1:
        e = 0
        while n % d == 0:
            n //= d
            e += 1
        if e > 0:
            fs += [(d, e)]
            root = ceil(sqrt(n))
        d += 1
    if n > 1:
        fs += [(n, 1)]
    return fs

# # Examples
# ```
# assert(factor_trial_division(73583) == [(73583, 1)])
# assert(factor_trial_division(73584) == [(2, 4), (3, 2), (7, 1), (73,
#     ↪ 1)])
# ```

```



Despite the naïveté of the algorithm, the trial division algorithm is often one of the very fastest when it comes to factoring small integers. In particular, if the integer to factor is a power of only two then trial division will outperform any of the other, more complex algorithms that we will soon cover. In fact, the main topic of this paper—the quadratic sieve algorithm—fails entirely not only for factoring primes but for factoring perfect powers in particular. (We will in fact be using trial division for one of the integral steps in some of the later algorithms.)

2.2 Fermat’s Method

Named after Baroque lawyer and mathematician Pierre de Fermat, Fermat’s method of factorization relies on a simple fact: that any odd integer is equal to the difference of squares.

As a simple visual proof of this fact, let N be any odd integer. By definition of N being odd, $N = 2m + 1$ for some integer b .

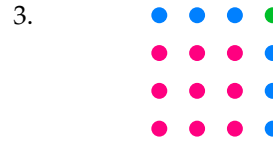
1. 

1. Draw a single  dot capped on either side by m  dots. The count of these dots is clearly N .

2. Bend the right side of ● dots down to form a square.
3. Fill the square with more ● dots.



Then clearly the area of the whole square minus the area of the inner square of only ● dots is equal to N (a difference of squares).



And if neither $a + b$ nor $a - b$ equal one or N , then they are *non-trivial* factors of N (the trivial factors of N being 1 and N itself).

With this fact in hand, let n be an odd composite integer; we wish to find c and d such that $n = cd$. (Note that we are only interested in odds since, if n were even, we simply factor out all copies of 2 until all that is left is supposedly an odd composite.) We know that for some a and b with $n = a^2 - b^2$ that the difference of squares is factorable into $(a + b)(a - b)$. This implies $c = a + b$ and $d = a - b$ with which we derive $a = (c + d)/2$ and $b = (c - d)/2$. We can verify:

$$n = \left(\frac{c+d}{2}\right)^2 - \left(\frac{c-d}{2}\right)^2 = \frac{(c^2 + 2cd - d^2) - (c^2 - 2cd + d^2)}{4} = cd.$$

Rewriting the difference of squares statement as $b^2 = a^2 - n$, and without loss of generality letting $a \geq \sqrt{n} > b$, the goal of Fermat's method in one way or another is to attempt to find an a satisfying the *strict* relationship $b = \sqrt{a^2 - n}$ where b is an integer. We can similarly implement this algorithm easily in Python but just checking all integers from $\lceil \sqrt{n} \rceil$ to n . Square roots are relatively cheap, but we do require a method of checking whether a root is an integer (we use $[b^2] \stackrel{?}{=} \lceil b \rceil^2$):

```
# Find two factors of an odd integer via Fermat's method
# @param n Odd integer to factor
# @return Two factors (possibly trivial)
def factor_fermat(n):
    for a in range(ceil(sqrt(n)), n + 1):
        b = sqrt(a**2 - n)
        if int(b**2) == ceil(b)**2:
            return(int(a + b), int(a - b))

# # Examples
# ```
# assert(factor_fermat(811) == (811, 1))
# assert(factor_fermat(813) == (271, 3))
# ```
```

As such there's not much improvement in terms of the complexity of this algorithm when compared to trial division; at worst case we have checked $\lceil \sqrt{n} \rceil$ integers when n is prime. Further more, we can only use this algorithm when n is odd, and we only get two particular factors for n . However, it is incredibly simple in this regard, and of course this simple implementation would only be part of a complete factoring algorithm which we could achieve by factoring 2 out from any even n to form an odd n , factoring every odd n via Fermat's method, and factoring each non-trivial resultant factor for as long as it takes until having completely factored n .

2.3 Congruences of Squares

In contrast, consider these *less-strict* congruence relation:

$$a^2 \equiv b^2 \pmod{n} \Rightarrow a^2 - b^2 = (a+b)(a-b) \equiv 0 \pmod{n}.$$

If n divides $(a+b)(a-b)$, then $\gcd(a+b, n)$ and $\gcd(a-b, n)$ certainly both yield factors of n . However, if $a \pm b \equiv 0 \pmod{n}$, then $\gcd(a \pm b, n) = n$ or 1, the trivial factors of n . So in addition to $a^2 \equiv b^2 \pmod{n}$ we add the constraint that $a \not\equiv \pm b \pmod{n}$, such that neither $a+b$ or $a-b$ are divisible by n . And if we find a pair a and b satisfying these two less-strict congruence relations

$$a^2 \equiv b^2 \pmod{n}, \tag{2.1}$$

$$a \not\equiv \pm b \pmod{n} \tag{2.2}$$

We have non-trivial factors of n .

Further more, the condition $a^2 \equiv b^2 \pmod{n}$ is that of b^2 being a *quadratic residue* modulo n . We say an integer b is a quadratic residue modulo n if there exists an x such that $x^2 \equiv b \pmod{n}$, and if no such x exists we call b a quadratic non-residue. If instead we take a prime p to be our modulus, then we have as a fact that of the integers from 0 to $p-1$ half are quadratic residues modulo p when p is odd (on the other hand, every integer is a quadratic residue modulo 2). Explicitly, we have the *Legendre symbol* of a modulo odd prime p :

$$\left(\frac{a}{p}\right) = a^{p-1/2} \pmod{p} \equiv \begin{cases} 0, & \text{if } p \text{ divides } a, \\ 1, & \text{if } a \text{ is a quadratic residue modulo } p, \text{ or} \\ -1, & \text{if } a \text{ is a quadratic non-residue modulo } p. \end{cases} \tag{2.3}$$

(These equivalent forms were expressed historically by Legendre and Euler.) Then *Jacobi symbol* $\left(\frac{a}{n}\right)$ generalizes the Legendre symbol for when instead the modulus is a composite n . Both the Legendre and Jacobi symbols have many useful properties, but in particular if n has the prime factorization $p_1^{e_1} p_2^{e_2} \cdots p_k^{e_k}$, then we have that the Jacobi symbol of a modulo n can be expressed as the product of powers of the Legendre symbols over each prime:

$$\left(\frac{a}{n}\right) = \left(\frac{a}{p_1}\right)^{e_1} \left(\frac{a}{p_2}\right)^{e_2} \cdots \left(\frac{a}{p_k}\right)^{e_k}. \tag{2.4}$$

Returning to the problem of factoring n , the question remains: *how do we find x and y pairs satisfying eq. (2.1) and eq. (2.2)?* The most naïve method would be to pick pairs x and y at random, hoping for a collision satisfying the conditions (but we will not exactly consider this an algorithm).

2.4 Dixon's Random Squares Algorithm

Consider that a composite integer is a perfect square only if every prime in its prime factorization is raised to a positive power (that is, every exponent has even parity). Factoring algorithms of the *random square* variety seek to find “small” composite integers such that their product yields a perfect square in this sense, to then solve the problem of congruences of squares. Dixon's algorithm/random squares method is one of such algorithms.

To do so practically, instead of factoring over the set of all the primes we select a subset \mathcal{B} of the t first primes. We call this our *factor base*, $\mathcal{B} = \{p_1, p_2, \dots, p_t\}$. We say that an integer n is p_t -smooth if the largest prime factor of n is less-than or equal-to the t^{th} prime, and thus smooth over our factor base of all the primes up to p_t .

And instead of looking for a and b pairs satisfying the (already weak) conditions from previously, by Dixon's method we look for integers a_i 's which satisfy the *even weaker still* relation $a_i^2 \equiv b_i \pmod{n}$ and where the b_i we calculate is smooth over our factor base (this step of checking whether b_i is p_t -smooth requires at least trial division and thus determining its prime factorization—we will need this later). In Dixon's method, the selection of these various a_i 's can be random, but considering that trial division may be expensive it is favorable to find them in a way where $b_i \pmod{n}$ ends up small. Two such ways, with integers $j \geq 0$ and $k \geq 1$, are $j + \lceil \sqrt{kn} \rceil$ and $\lfloor \sqrt{kn} \rfloor$; the former produces integers $a_i^2 \pmod{n}$ that tend small, and the later produces integers $-a_i^2 \pmod{n}$ that also tend to be small. (One thing to note is that in the case of the later we need to add $p_0 = -1$ to the factor base to be able to factor anything in the form $-a_i^2$; ultimately this doesn't change anything about the algorithm aside that it can greatly improve the speed at which we find pairs since we have twice as many potential small a_i 's.)

But once we have found pairs (a_i, b_i) satisfying $a_i^2 \equiv b_i \pmod{n}$ and b_i is p_t -smooth, *what do we do with them?* Recall the consideration (three paragraphs back) that an integer is a perfect square if and only if its prime power factorization is of even powers of primes. For each of the (a_i, b_i) pairs we express b_i as a product of powers of the primes in \mathcal{B} our factor base and record these exponents as vectors in both decimal and over the field \mathbb{Z}_2 :

$$b_i = p_1^{e_{i1}} \cdot p_2^{e_{i2}} \cdots p_t^{e_{it}} \quad \Rightarrow \quad \mathbf{e}_i = \langle e_{i1}, e_{i2}, \dots, e_{it} \rangle \quad \Rightarrow \quad \mathbf{v}_i = \langle v_{i1}, v_{i2}, \dots, v_{it} \rangle \quad (2.5)$$

For each $p_j \in \mathcal{B}$, and where $v_{ij} \equiv e_{ij} \pmod{2}$ (i.e. the parity of e_{ij}). Recall that we know the prime power factorization of b_i from checking that b_i is p_t -smooth via trial division. If the sum of a subset of these \mathbf{v}_i equals the zero vector, *then we have found a congruence of squares satisfying $a^2 \equiv b^2 \pmod{n}$* , just with a few extra steps left to construct a from the a_i 's and b from the b_i 's. In terms of linear algebra, we have found a linearly dependent subset of the \mathbf{v}_i vectors. This clues us into quantifying how many (a_i, b_i) pairs we need to *guarantee* that we actually have enough such pairs:

Our vectors have t columns, one for every prime in \mathcal{B} the factor base; if we form a matrix of the vectors, then the rank of this matrix is at most $|\mathcal{B}|$ the number of columns; lastly, if the number of rows of the matrix is equal to one plus its rank (the number of columns), then we *guarantee* that the row space contains linear dependencies (since fundamentally the row rank and column rank of a matrix are always equal, and at “worst” the column space is entirely linearly dependent with rank $|\mathcal{B}|$).

Thus we need *at least* $|\mathcal{B}| + 1$ (a_i, b_i) pairs to guarantee that we can construct a congruence of squares.

Constructing a and b is actually fairly straight forward: let $I \subseteq \{1, 2, 3, \dots, t\}$ be the set of indices i for which we found $\sum_{i \in I} v_i = \mathbf{0}$. Fix b^2 to be:

$$b^2 = \prod_{i \in I} b_i = \prod_{j=1}^t p_j^{\epsilon_j} \equiv \left(\prod_{i \in I} a_i \right)^2 = a^2 \pmod{n}. \quad (2.6)$$

Where because each $\epsilon_j = \prod_{i \in I} e_{ij}$ we have positive parity we know b^2 to be a perfect square. Thus we can construct b by simply dividing each ϵ_j by two:

$$b = \left(\prod_{i \in I} b_i \right)^{1/2} = \prod_{j=1}^t p_j^{\epsilon_j/2}. \quad (2.7)$$

It suffices to calculate $a = \prod_{i \in I} a_i$ since we already have that a^2 is clearly a perfect square in the satisfied congruence relation [eq. \(2.6\)](#).

Despite all the work we have done, if we find that $a \equiv \pm b \pmod{n}$ then the pair (a, b) has unfortunately failed to yield anything but the trivial factors of n . As for our avenues of recourse, we can:

1. Check for other linearly dependent subsets I and their corresponding pairs (a, b) . If every subset failed, then
2. We can search for more (a_i, b_i) pairs (we call this increasing the search interval), perform the previous actions all over again but with these new pairs, then if that fails
3. Start over with an expanded factor base.

2.5 The Quadratic Sieve Algorithm

Last but not least we discuss the QS algorithm, but despite all this lead-up there is truly not that much more to it in comparison to Dixon’s. The QS algorithm functions almost identically to Dixon’s method, except for the way in which potential a_i ’s are selected: first by a simple translation, and second in actual “sieving.”

Letting n be the integer to factor and $m = \lfloor \sqrt{n} \rfloor$, consider the polynomial $f(x) = (x + m)^2 - n$. In Dixon’s algorithm we selected a_i ’s either randomly or using a slight

heuristic. Instead, in the QS algorithm we select $a_i = (x + m)$, calculate a corresponding $b_i = (x + m)^2 - n$ and test whether b_i is p_t -smooth (just like in Dixon's). Note the relationship between a_i and b_i pair:

$$a_i^2 = (x + m)^2 = (x + m)^2 - n + n \equiv (x + m)^2 - n = b_i \pmod{n}. \quad (2.8)$$

Identical to the congruence relation from Dixon's. Recall that we want the b_i 's to factor entirely over our chosen factor base \mathcal{B} (that is to be p_t -smooth), but suppose b_i is divisible by some prime p in general. Then p also divides $(x + m)^2 - n$ (from our definition of b_i), which gives us

$$(x + m)^2 - n \equiv 0 \pmod{p} \Rightarrow (x + m)^2 \equiv n \pmod{p}. \quad (2.9)$$

Clearly $(x + m)^2$ is a square, meaning that n is a quadratic residue (modulo p)! Therefore we can narrow down the factor base to contain only those primes p_j such that n is a quadratic residue modulo p_j .

Chapter 3

Implementation

I decided to designate the implementation discussion to the Rust implementation as opposed to the Python implementation. Not because it is more complete than the Python implementation (on the contrary even), but because I find it vastly more interesting. Still, I wrote the Python implementation first as a sanity check (it is definitely easier and faster to bang-out an implementation in Python than Rust) and its entirety can be found here: github.com/nilsso/quadratic-sieve-rs/python-version/.

For those who wish to peruse it, begin in the `qs-sieving.py` module, which can then lead you to the other modules and how they are used (like the Tonelli-Shanks `sqrt_mod` algorithm, which leads to the Legendre symbol implementation, the `is_quadratic_prime` function and the `legendre_primes` iterator in `legendre.py`).

With the fundamental algorithm completed in Python, I made my move to Rust. This section is primarily a retrospective on my experience in implementing (*or at least attempting to implement*) the QS algorithm in Rust. I intend for it to be useful to anyone with an interest in Rust, as well as useful to myself in both reviewing my own knowledge of the language and providing a figurative map of my thoughts on the design of my project's implementation. I go over the Rust language in general, how its features have been very useful for implementing mathematical structures, and weave in how I used these features in my own QS implementation. If heavy use of programming jargon is not your thing then this section may not be very interesting. That being said, the full source code of the Rust portion of this project can be found here: github.com/nilsso/quadratic-sieve-rs/.

3.1 Implementing in Rust

As far as fundamental paradigms go, Rust as a programming language has much in common with more classically object oriented languages like C++ or Java, but also a lot in common with purely function languages like Haskell. Similarities with the later are in Rust's

sophisticated *trait* feature, analogous to Haskell's *type classes*, with which Rust accomplishes polymorphism. In C++, polymorphism is achieved through hierarchies of classes and abstract classes (i.e. interfaces) (e.g. if a Square derives from Shape, then a function which accepts a Shape can accept a Square). Although Rust encapsulates data and functionality through structs (analogous to C++ classes) and their instantiations, polymorphism in Rust is achieved *only* through its traits (e.g. if Shape is a trait and Square implements the Square trait, then a function that requires its argument implement Shape can accept a Square).

The most substantial difference in polymorphism is that there is *no such thing* as inheritance for Rust structs like there is for C++ classes. (e.g. in C++, Square as a class is a subclass of an abstract class Shape; in Rust, Square *implements* Shape, where Square is a struct and Shape is a trait. Further more, *if* Shape *was* a struct *instead of* a trait, Square could not inherit any of Shape's functionality and neither would share any functionality—as far as the compiler is concerned—without them both implementing some other trait.) This is a major paradigm shift for anyone moving from C++ to Rust, as inheritance based polymorphism is more-or-less the only flavor of polymorphism that students of computer science/programming are likely to have learned (as was the case for me especially), and one might believe on a first glance that if inheritance is not a feature then Rust structs cannot share functionality without copying entire swathes of code.

Instead, Rust traits fill this purpose: a trait can define default functionality for which any struct which implements the trait inherits. For example, the code in [listing 1](#) defines a trait Pow representing the capability of the implementing type to be “raised to a power,” for which any type can implement. (A note on the language: an `impl` block *implements* functionality, i.e. one or several functions, onto an existing type. These can be *methods* for instances of the type, like `a.foo()`, or *functions* belonging to the type itself, like `Foo::foo()`.)

```
trait Pow {
    /// Raise this value to an exponent.
    /// (Must be implemented)
    fn pow(&self, e: u32) → Self;

    /// Square this value.
    /// (Implemented by default, and does *not* need to be re-implemented)
    fn squared(&self) → Self {
        self.pow(2)
    }
}

// Although we don't implement the `squared` method for `i32` here,
// it gets the default implementation from the `Pow` trait.
impl Pow for i32 {
    fn pow(&self, e: u32) → i32 {
        <i32>::pow(e)
    }
    // fn squared... is inherited.
}
```

Listing 1: A trait with a default implementation.

The fundamental advantage of traits used in this way is that if we then write a function in which we need the functionality of `pow`, then we can require that whatever type(s) the function works over, and/or takes as parameters, *implement the `Pow` trait*; furthermore, anyone is allowed to implement `Pow` trait, and thus use our function with their own `Pow` implementing type. In other words, *polymorphism in Rust is achieved by abstracting over the functionality that types have, and not the types themselves in particular*. We call these functionality requirements on types *trait bounds*.

Another advantage of trait bounds is that they allow us to add functionality to a type *incrementally* by requiring increasingly stricter trait bounds for more functionality dependent operations. I bring up an example that I will use continuously from now on: my implementation of $M \times N$ matrices, `Matrix`. At its most generic, a matrix can be just a two-dimensional storage location in which we could potentially store anything. As such, in [listing 2](#) I define `Matrix` without any bounds on what it is allowed to store aside from the type needing to be clonable. (Heed not the syntax `const M` and `const N`, but know that they control the size of the matrix; we will cover this syntax in a few paragraphs from here)

```
#[derive(Clone)]
pub struct Matrix<T, const M: usize, const N: usize>
where
    T: Clone,
{
    pub rows: [[T; N]; M], // M arrays of T arrays of length N
}
```

Listing 2: The verbatim definition of my own matrix type.

In this snippet, `T` is called a generic parameter, and it is given the trait bound `Clone`; thus we are allowed to make a matrix of type `T` for any type `T` that is clonable. Personally, I find something inherently mathematical about the way in which we can restrict generic types in Rust. This leads me to a particular implementation ([listing 3](#)) for my `Matrix` struct: the zero matrix.

```
impl<T, const M: usize, const N: usize> Matrix<T, M, N>
where
    T: Zero<Element = T> + Clone,
{
    pub fn zeroes() → Self {
        // T::ZERO is the constant value that represents
        // zero for the generic type parameter T.
        let rows = array_init(|_| array_init(|_| T::ZERO));
        Self::from_array(rows)
    }
}
```

Listing 3: Implementation to construct a $M \times N$ all-zero matrix.

In this snippet we define the construction of a zero matrix of specified size $M \times N$. This begs the question: *how do we know that the type T , for which the matrix is over, has a representation of zero at all?* For the generic type T , the only way then for the compiler to know that T has zero is to impose an additional trait bound, in this case a trait of my own called `Zero`. The only thing that this trait requires is that an implementing type provide a constant representation of zero, and with this we have a value with which to fill a zero matrix. Additionally, note the `Element = T` part of the trait bound: this requires that not only does T need to have a zero element, but that the zero element *must be of the same type as T itself*. This allows for a distinction when, for example, something like a set of numbers has a zero element, but the zero element is not itself a set (think of the quotient group \mathbb{Z}_n having the congruence class $[0]_n$ as its zero element). (See [listing 4](#) for the implementation of `Zero` on `i32` integers and on my own `CongruenceClass` struct; I cover the later in more detail next).

```
pub trait Zero {
    type Element;

    const ZERO: Self::Element;
}

pub trait One {
    type Element;

    const ONE: Self::Element;
}

impl Zero for i32 {
    type Element = i32;

    const ZERO: i32 = 0;
}

impl<const M: u32> Zero for QuotientGroup<M> {
    type Element = CongruenceClass<M>;

    const ZERO: CongruenceClass<M> = CongruenceClass<M>(0);
}
```

Listing 4: Definition of the zero and one traits, and the implementation of zero for the 32-bit integer type, and for my own quotient group type.

A highly experimental feature of Rust that I only discovered since having started this project is `const generics`. This feature has proven to be incredibly useful in the implementation of various mathematical structures, in particular congruence classes (see [listing 5](#)) and matrices.

```
/// Congruence class with modulo M.
```

```

///
/// # Examples
/// We can add congruence classes if they have the same modulus:
/// ```
/// let a = CongruenceClass::<5>::new(1); // 1 modulo 5
/// let b = CongruenceClass::<5>::new(4); // 4 modulo 5
/// 1 + 4 is congruent to 0 modulo 5:
/// assert_eq!(a + b, CongruenceClass::<5>::new(0));
/// ```
/// But not if they have different modulus:
/// ```skip
/// let a = CongruenceClass::<5>::new(1); // 1 modulo 5
/// let b = CongruenceClass::<6>::new(5); // 5 modulo 6
/// a + b; // fails to compile, since 5 ≠ 6
/// ```
pub struct CongruenceClass<const M: u32>(pub u32);

impl<const M: u32> CongruenceClass<M> {
    pub fn new(x: u32) → Self {
        Self(x % M)
    }
}

// `Add` implies `Add<CongruenceClass<M>>`, and
// `Self` refers to `CongruenceClass<M>`
impl<const M: u32> Add for CongruenceClass<M> {
    type Output = Self;

    fn add(self, rhs: Self) → Self {
        Self((self.0 + rhs.0) % M)
    }
}

```

Listing 5: Excerpt of my implementation of congruence classes in Rust, using traits and the const generic feature.

What const generics do is bake a constant value (though currently only primitive types like `u32` are supported) directly into a trait or struct. This allows the compiler to know at compile time (or more specifically the type checker and type check time) whether or not, say, two instances of a struct are compatible with one-another based on the constants that are baked into their types; because although a may have to instances of `CongruenceClass`, their types are not fully qualified without the constant modulus.

Another instance in which this feature was useful was in the defining of my matrix struct. What is amazing about the const generic feature in this case is that the validity checking of operations like matrix addition and multiplication no longer need to take place a runtime, and instead are *completely validated at compile time!* (See the implementation of addition and multiplication in [listing 6](#) below.)

```

// Implement same size matrix addition
impl<T, const M: usize, const N: usize> Add for Matrix<T, M, N> {
    type Output = Self;

    fn add(self, rhs: Self) → Self { /* ... */ }
}

// Implement MxN and NxP → MxP matrix multiplication
//
// With the use of const generics we can ensure at compile time that
// → matrices
// that we multiply must have the same number of columns on the left as
// → rows
// on the right.
impl<T, const M: usize, const N: usize, const P: usize> Mul for Matrix<T,
    → M, N>
where
    T: Add<Output = T> + Mul<Output = T> + Copy,
{
    type Output = Matrix<T, M, P>;

    fn mul(self, rhs: Matrix<T, N, P>) → Matrix<T, M, P> { /* ... */ }
}

```

Listing 6: Excerpt of my implementation of matrices in Rust, using traits and the const generic feature.

Checking validity before the program is ever ran allows us to guarantee the success of such operations at runtime, eliminating the perhaps drastic amount of time spent checking whether matrices have compatible shapes at runtime otherwise.

Now, because `Matrix` was defined generically, it was trivial to use it and `CongruenceClass` in conjunction to solve for the linear dependent subsets of the binary exponent vectors. If I know B the size of the factor base and have acquired $B + 1$ smooth points, then I can map the exponent vectors of the prime factorizations of the points into a matrix over \mathbb{Z}_2 ; in Rust the type of this matrix looks like `Matrix< CongruenceClass<2>, B+1, B>`. Having **implemented the row-reduction of a matrix into Echelon form** for any generic type `T` (achievable given a variety of trait bounds), then nothing additional need be implemented for solving for the spanning vectors of the exponent matrix left nullspace over the particular type \mathbb{Z}_2 (that is, `CongruenceClass<2>` in Rust). And thus the problem of solving for the linearly dependent exponent vectors was solved. And, up to this point, with these tools (and with implementing a few other helper algorithms) I was additionally able to implement the quadratic sieving algorithm itself (that is, finding $B + 1$ numbers that are smooth over the factor base, guaranteeing linear dependence between their binary exponent vectors).

However, this is where the story for the time being turns sour, as I am currently *not* able to complete my implementation of the QS algorithm based on the fact that my matrix struct requires knowing the number of rows M and columns N *at compile time*. Depend-

ing on the number to factor n , there is no way for my program to deterministically know at compile time how large the factor base needs to be (controlling the number of columns of the matrix) nor how many smooth numbers it finds (controlling the number of rows). And given that `Matrix` requires that its dimensions M and N be known at compile time, I currently cannot construct the necessary matrices from the exponent vectors. A solution will be found eventually (for instance I should be able to change the factor base length argument into a const generic of the function itself, and proceed), but for now this implementation is dead in the water, at this crucial step.

As mentioned in the beginning of this section, the full source code can be found within:

- github.com/nilsso/quadratic-sieve-rs/

The Rust implementation is not nearly as well documented as the Python version yet, but the structure of the Cargo (Rust's default project manager and tooling) project looks as follows:

- **quadratic_sieve.rs** contains what is finished of the main algorithm, including the **sieve itself, an implementation** of the Tonelli-Shanks algorithm for calculating square-roots modulo a prime p , and various other helper functions and structs.
- **tests/** contains a few unit tests for sanity checking on the congruence class and matrix structs.
- **matrix.rs** contains the `Matrix` definition and its implementations.
- **congruence_class.rs** contains the `CongruenceClass` definition and its implementations.
- **identity.rs** contains the `Zero` and `One` traits, and their primitive type implementations.
- **integers.rs** contains various traits for representing common functionality between integer (or integer-like) types, and as part includes implementations of **the extended Euclidean algorithm** (for calculating GCD, LCM and modular inverses), **Stein's binary GCD algorithm** (currently used for calculating the GCD of unsigned integers), and **repeated squaring modulo m** , among others.
- As well as a few other miscellaneous bits of code, either for fun like with **quotient groups, complex numbers** and the **Conjugate trait** (of which `Matrix` implements as long as `T` implements `Conjugate`; i.e. integers and complex numbers).