

Quadratic Sieve Algorithm

a.k.a. the Second Fastest Integer Factoring Algorithm in the West
(and trying to get it to work)

Nils Olsson

Spring Semester, 2021

Contents

1	Introduction	2
2	Square Factoring	3
3	The Quadratic Sieve	4
4	Implementation	5
4.1	Implementing in Rust	5
4.1.1	Overview of Rust	5

Abstract

This report serves as a summary of the history and mathematical theory behind square factoring functions, as an loose introduction to the Rust programming language, and as a retrospective on my own implementing of the quadratic sieve (QS) algorithm.

Chapter 1

Introduction

The factorization of integers has been an especially well studied problem, and not just for centuries but in fact millennia. Greek mathematicians studied the problem, even proving the fundamental theorem of arithmetic: that every integer has a unique prime factorization. As a corollary the Greeks had also extensively studied the prime integers, with the sieve of Eratosthenes being one of the most important algorithms for generating primes ever invented. The so called father of geometry Euclid of Alexandria (whose now self-titled “Euclid’s lemma” was foundational for proving integer factorization uniqueness) described in his *Elements* the Euclidean algorithm for calculating the greatest common divisor of two integers. These ancient algorithms, and many others, comprise the fundamental building blocks for which all modern integer factorization algorithms are made.

Chapter 2

Square Factoring

Chapter 3

The Quadratic Sieve

Chapter 4

Implementation

I decided to designate the implementation discussion to the Rust implementation as opposed to the Python implementation. Not because it is more complete than the Python implementation (on the contrary even), but because I find it vastly more interesting. Still, I wrote the Python implementation first as a sanity check (it is definitely easier and faster to bang-out an implementation in Python than Rust) and its entirety can be found here: github.com/nilsso/quadratic-sieve-rs/python-version/.

For those who wish to peruse it, begin in the [qs-sieving.py](#) module, which can then lead you to the other modules and how they are used (like the Tonelli-Shanks [sqr_mod](#) algorithm, which leads to the Legendre symbol, `is_quadratic_prime` function and `legendre_primes` iterator in [legendre.py](#)).

4.1 Implementing in Rust

With the fundamental algorithm completed in Python, I made the move to re-implementing in the Rust language. This section is primarily a retrospective on my experience in implementing (*or at least attempting to implement*) the QS algorithm in Rust. I intend for it to be useful to anyone with an interest in Rust, as well as useful to myself in both reviewing my own knowledge of the language and providing a figurative map of my thoughts on the design of my project's implementation. I go over the Rust language in general, how its features have been very useful for implementing mathematical structures, and weave in how I used these features in my own QS implementation. If heavy use of programming jargon is not your thing then this section may not be very interesting. That being said, the full source code of the Rust portion of this project can be found here

- github.com/nilsso/quadratic-sieve-rs/

4.1.1 Overview of Rust

As far as fundamental paradigms go, Rust as a programming language has much in common with more classically object oriented languages like C++ or Java, but also a lot in common with purely function languages like Haskell. Similarities with the later are in Rust's sophisticated *trait* feature, analogous to Haskell's *type classes*, with which Rust accomplishes polymorphism. In C++, polymorphism is achieved through hierarchies of `classes` and `abstract classes` (i.e. `interfaces`) (e.g. if a `Square` derives from `Shape`, then a function which accepts a `Shape` can accept a `Square`). Although Rust encapsulates data and func-

tionality through structs (analogous to C++ classes) and their instantiations, polymorphism in Rust is achieved *only* through its traits (e.g. if Shape is a trait and Square implements the Square trait, then a function that requires its argument implement Shape can accept a Square).

The most substantial difference in polymorphism is that there is *no such thing* as inheritance for Rust structs like there is for C++ classes. (e.g. in C++, Square as a class is a subclass of an abstract class Shape; in Rust, Square *implements* Shape, where Square is a struct and Shape is a trait. Further more, *if* Shape *was* a struct *instead of* a trait, Square could not inherit any of Shape’s functionality and neither would share any functionality—as far as the compiler is concerned—without them both implementing some other trait.) This is a major paradigm shift for anyone moving from C++ to Rust, as inheritance based polymorphism is more-or-less the only flavor of polymorphism that students of computer science/programming are likely to have learned (as was the case for me especially), and one might believe on a first glance that if inheritance is not a feature then Rust structs cannot share functionality without copying entire swathes of code.

Instead, Rust traits fill this purpose: a trait can define default functionality for which any struct which implements the trait inherits. For example, the code in [listing 1](#) defines a trait Pow representing the capability of the implementing type to be “raised to a power,” for which any type can implement. (A note on the language: an impl block *implements* functionality, i.e. one or several functions, onto an existing type. These can be *methods* for instances of the type, like `a.foo()`, or *functions* belonging to the type itself, like `Foo::foo()`.)

```
trait Pow {
    /// Raise this value to an exponent.
    /// (Must be implemented)
    fn pow(&self, e: u32) → Self;

    /// Square this value.
    /// (Implemented by default, and does *not* need to be re-implemented)
    fn squared(&self) → Self {
        self.pow(2)
    }
}

// Although we don't implement the `squared` method for `i32` here,
// it gets the default implementation from the `Pow` trait.
impl Pow for i32 {
    fn pow(&self, e: u32) → i32 {
        <i32>::pow(e)
    }
    // fn squared... is inherited.
}
```

Listing 1: A trait with a default implementation.

The fundamental advantage of traits used in this way is that if we then write a function in which we need the functionality of pow, then we can require that whatever type(s) the function works over, and/or takes as parameters, *implement the Pow trait*; furthermore, anyone is allowed to implement Pow trait, and thus use our function with their own Pow implementing type. In other words, *polymorphism in Rust is achieved by abstracting over the functionality that types have, and not the types themselves in particular*. We call these functionality requirements on types *trait bounds*.

Another advantage of trait bounds is that they allow us to add functionality to a type *incrementally* by requiring increasingly stricter trait bounds for more functionality dependent operations. I bring up an example that I will use continuously from now on: my implementation of $M \times N$ matrices, `Matrix`. At its

most generic, a matrix can be just a two-dimensional storage location in which we could potentially store anything. As such, in [listing 2](#) I define `Matrix` without any bounds on what it is allowed to store aside from the type needing to be clonable. (Heed not the syntax `const M: usize` and `const N: usize`, but know that they control the size of the matrix; we will cover this syntax in a few paragraphs from here)

```
#[derive(Clone)]
pub struct Matrix<T, const M: usize, const N: usize>
where
    T: Clone,
{
    pub rows: [[T; N]; M], // M arrays of T arrays of length N
}
```

Listing 2: The verbatim definition of my own matrix type.

In this snippet, `T` is called a generic parameter, and it is given the trait bound `Clone`; thus we are allowed to make a matrix of type `T` for any type `T` that is clonable. Personally, I find something inherently mathematical about the way in which we can restrict generic types in Rust. This leads me to a particular implementation ([listing 3](#)) for my `Matrix` struct: the zero matrix.

```
impl<T, const M: usize, const N: usize> Matrix<T, M, N>
where
    T: Zero<Element = T> + Clone,
{
    pub fn zeroes() → Self {
        // T::ZERO is the constant value that represents
        // zero for the generic type parameter T.
        let rows = array_init(|_| array_init(|_| T::ZERO));
        Self::from_array(rows)
    }
}
```

Listing 3: Implementation to construct a $M \times N$ all-zero matrix.

In this snippet we define the construction of a zero matrix of specified size $M \times N$. This begs the question: *how do we know that the type `T`, for which the matrix is over, has a representation of zero at all?* For the generic type `T`, the only way then for the compiler to know that `T` has zero is to impose an additional trait bound, in this case a trait of my own called `Zero`. The only thing that this trait requires is that an implementing type provide a constant representation of zero, and with this we have a value with which to fill a zero matrix. Additionally, note the `Element = T` part of the trait bound: this requires that not only does `T` need to have a zero element, but that the zero element *must be of the same type as `T` itself*. This allows for a distinction when, for example, something like a set of numbers has a zero element, but the zero element is not itself a set (think of the quotient group \mathbb{Z}_n having the congruence class $[0]_n$ as its zero element). (See [listing 4](#) for the implementation of `Zero` on `i32` integers and on my own `CongruenceClass` struct; I cover the later in more detail next).

```
pub trait Zero {
    type Element;

    const ZERO: Self::Element;
}

pub trait One {
    type Element;
```



```

    const ONE: Self::Element;
}

impl Zero for i32 {
    type Element = i32;

    const ZERO: i32 = 0;
}

impl<const M: u32> Zero for QuotientGroup<M> {
    type Element = CongruenceClass<M>;

    const ZERO: CongruenceClass<M> = CongruenceClass<M>(0);
}

```

Listing 4: Definition of the zero and one traits, and the implementation of zero for the 32-bit integer type, and for my own quotient group type.

A highly experimental feature of Rust that I only discovered since having started this project is `const` generics. This feature has proven to be incredibly useful in the implementation of various mathematical structures, in particular congruence classes (see [listing 5](#)) and matrices.

```

/// Congruence class with modulo M.
///
/// # Examples
/// We can add congruence classes if they have the same modulus:
/// ```
/// let a = CongruenceClass::<5>::new(1); // 1 modulo 5
/// let b = CongruenceClass::<5>::new(4); // 4 modulo 5
/// 1 + 4 is congruent to 0 modulo 5:
/// assert_eq!(a + b, CongruenceClass::<5>::new(0));
/// ```
/// But not if they have different modulus:
/// ```skip
/// let a = CongruenceClass::<5>::new(1); // 1 modulo 5
/// let b = CongruenceClass::<6>::new(5); // 5 modulo 6
/// a + b; // fails to compile, since 5 ≠ 6
/// ```
pub struct CongruenceClass<const M: u32>(pub u32);

impl<const M: u32> CongruenceClass<M> {
    pub fn new(x: u32) → Self {
        Self(x % M)
    }
}

// `Add` implies `Add<CongruenceClass<M>>`, and
// `Self` refers to `CongruenceClass<M>`
impl<const M: u32> Add for CongruenceClass<M> {
    type Output = Self;

    fn add(self, rhs: Self) → Self {
        Self((self.0 + rhs.0) % M)
    }
}

```

Listing 5: Excerpt of my implementation of congruence classes in Rust, using traits and the const generic feature.

What const generics do is bake a constant value (though currently only primitive types like u32 are supported) directly into a trait or struct. This allows the compiler to know at compile time (or more specifically the type checker and type check time) whether or not, say, two instances of a struct are compatible with one-another based on the constants that are baked into their types; because although a may have to instances of CongruenceClass, their types are not fully qualified without the constant modulus.

Another instance in which this feature was useful was in the defining of my matrix struct. What is amazing about the const generic feature in this case is that the validity checking of operations like matrix addition and multiplication no longer need to take place a runtime, and instead are *completely validated at compile time!* (See the implementation of addition and multiplication in [listing 6](#) below.)

```
// Implement same size matrix addition
impl<T, const M: usize, const N: usize> Add for Matrix<T, M, N> {
    type Output = Self;

    fn add(self, rhs: Self) → Self { /* ... */ }
}

// Implement MxN and NxP → MxP matrix multiplication
//
// With the use of const generics we can ensure at compile time that matrices
// that we multiply must have the same number of columns on the left as rows
// on the right.
impl<T, const M: usize, const N: usize, const P: usize> Mul for Matrix<T, M, N>
where
    T: Add<Output = T> + Mul<Output = T> + Copy,
{
    type Output = Matrix<T, M, P>;

    fn mul(self, rhs: Matrix<T, N, P>) → Matrix<T, M, P> { /* ... */ }
}
```

Listing 6: Excerpt of my implementation of matrices in Rust, using traits and the const generic feature.

Checking validity before the program is ever ran allows us to guarantee the success of such operations at runtime, eliminating the perhaps drastic amount of time spent checking whether matrices have compatible shapes at runtime otherwise.

Now, because Matrix was defined generically, it was trivial to use it and CongruenceClass in conjunction to solve for the linear dependent subsets of the binary exponent vectors. If I know B the size of the factor base and have acquired $B + 1$ smooth points, then I can map the exponent vectors of the prime factorizations of the points into a matrix over \mathbb{Z}_2 ; in Rust the type of this matrix looks like `Matrix<CongruenceClass<2>, B+1, B>`. Having **implemented the row-reduction of a matrix into Echelon form** for any generic type T (achievable given a variety of trait bounds), then nothing additional need be implemented for solving for the spanning vectors of the exponent matrix left nullspace over the particular type \mathbb{Z}_2 (that is, `CongruenceClass<2>` in Rust). And thus the problem of solving for the linearly dependent exponent vectors was solved. And, up to this point, with these tools (and with implementing a few other helper algorithms) I was additionally able to implement the quadratic sieving algorithm itself (that is, finding $B + 1$ numbers that are smooth over the factor base, guaranteeing linear dependence between their binary exponent vectors).

However, this is where the story for the time being turns sour, as I am currently *not* able to complete my implementation of the QS algorithm based on the fact that my `matrix struct` requires knowing the number of rows M and columns N *at compile time*. Depending on the number to factor n , there is no way for my program to deterministically know at compile time how large the factor base needs to be (controlling the number of columns of the matrix) nor how many smooth numbers it finds (controlling the number of rows). And given that `Matrix` requires that its dimensions M and N be known at compile time, I currently cannot construct the necessary matrices from the exponent vectors. A solution will be found eventually (for instance I should be able to change the factor base length argument into a `const generic` of the function itself, and proceed), but for now this implementation is dead in the water, at this crucial step.

As mentioned in the beginning of this section, the full source code can be found within:

- github.com/nilsso/quadratic-sieve-rs/

The Rust implementation is not nearly as well documented as the Python version yet, but the structure of the Cargo (Rust's default project manager and tooling) project looks as follows:

- **`quadratic_sieve.rs`** contains what is finished of the main algorithm, including the **sieve itself**, an **implementation** of the Tonelli-Shanks algorithm for calculating square-roots modulo a prime p , and various other helper functions and structs.
- **`tests/`** contains a few unit tests for sanity checking on the congruence class and matrix structs.
- **`matrix.rs`** contains the `Matrix` definition and its implementations.
- **`congruence_class.rs`** contains the `CongruenceClass` definition and its implementations.
- **`identity.rs`** contains the `Zero` and `One` traits, and their primitive type implementations.
- **`integers.rs`** contains various traits for representing common functionality between integer (or integer-like) types, and as part includes implementations of **the extended Euclidean algorithm** (for calculating GCD, LCM and modular inverses), **Stein's binary GCD algorithm** (currently used for calculating the GCD of unsigned integers), and **repeated squaring modulo m** , among others.
- As well as a few other miscellaneous bits of code, either for fun like with **quotient groups**, **complex numbers** and the **Conjugate trait** (of which `Matrix` implements as long as `T` implements `Conjugate`; i.e. integers and complex numbers).